

Computer Vision
Mini-project 3
CS 670, Fall 2019

Name: Kunjal Panchal
Student ID: 32126469
Email: kpanchal@umass.edu

October 23, 2019

Contents

1	Problem 1: Image Denoising	3
1.1	(a) Gaussian Filtering	3
1.2	(b) Median Filtering	6
1.3	(c) Non-Local Means Filtering	8
1.4	Results and Conclusions	11
2	Problem 2: Texture Synthesis	14
2.1	Objective	14
2.2	(a) Random Tiling	15
2.2.1	Results	15
2.3	(b) Non-parametric Sampling	17
2.3.1	Outputs	17
2.3.2	Effect of window size on runtime and on synthesis quality	17
3	Extensions for Extra Credit	21
3.1	Extension 1: Block Matching 3D (BM3D)	21
3.2	Extension 2: Image Quilting	24
A		
	Matlab code for Problem 1: Image Denoising	28
A.1	Implementation of <code>evalDenoising.m</code>	28
A.2	Implementation of Non-Local Means Filters	34
B		
	Matlab code for Problem 2: Texture Synthesis	38
B.1	Implementation of <code>evalTextureSynth.m</code>	38
B.2	Implementation of <code>synthRandomPatch.m</code>	39
B.3	Implementation of <code>synthEfrosLeung.m</code>	40
C		
	Matlab code for Extension for Extra Credit	45
C.1	Implementation of <code>synthImageQuilting.m</code>	45

1 Problem 1: Image Denoising

Let's take a look at some terms useful for understanding this problem:

- **IMAGE DENOISING** Images taken with both digital cameras and conventional film cameras will pick up noise from a variety of sources. Further use of these images will often require that the noise be (partially) removed.
- **GAUSSIAN NOISE** Variations in intensity drawn from a Gaussian normal distribution. [1]
- **SALT AND PEPPER NOISE** Contains random occurrences of black and white pixels.
- **LINEAR SMOOTHING FILTERS** One method to remove noise is by convolving the original image with a mask that represents a low-pass filter or smoothing operation. For example, the Gaussian mask comprises elements determined by a Gaussian function. [2]
- **NON LINEAR FILTERS** A median filter is an example of a non-linear filter and, if properly designed, is very good at preserving image detail.
- **NON LOCAL MEANS** Another approach for removing noise is based on non-local averaging of all the pixels in an image. In particular, the amount of weighting for a pixel is based on the degree of similarity between a small patch centered on that pixel and the small patch centered on the pixel being de-noised. [2]

1.1 (a) Gaussian Filtering

The idea of Gaussian smoothing is to use this 2-D distribution as a 'point-spread' function, and this is achieved by convolution.

Since the image is stored as a collection of discrete pixels we need to produce a discrete approximation to the Gaussian function before we can perform the convolution.

In theory, the Gaussian distribution is non-zero everywhere, which would require an infinitely large convolution kernel, but in practice it is effectively zero more than about three standard deviations from the mean,

and so we can truncate the kernel at this point.

Figure 1 shows a suitable integer-valued convolution kernel that approximates a Gaussian with a σ of 1.0.

$$\frac{1}{273}$$

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

Figure 1: Discrete approximation to Gaussian function with $\sigma=1.0$

We convolute all noisy images with gaussian filters first, we also experiment with different σ standard deviations and find out which value gives the optimal results.

NOTE: Matlab support says that the recommended replacement functions, "imgaussfilt" and "imgaussfilt3", are "generally faster and more advanced" with comparison to "fspecial('gaussian',hsize,sigma)". Also they point out that the documentation for "imgaussfilt" states that the "A" matrix input can be of any dimension, so it should work for N-D images.

Hence, we use "imgaussfilt" instead of combination of "fspecial" (to create a gaussian filter kernel) and "imfilter" (to convolute it with noisy image and get a denoised image).

This is shown in **line 7 and 8** of code snippet 1.1.

Matlab Implementation of Gaussian Filters

```

1 %% Denoising algorithm (Gaussian filtering)
2
3 count = 1;
4 % Experiment with different values of standard
  deviation

```

Image Name	Optimal Standard Deviation σ	Least Error - SSD
saturn-noise1g.png	1.09	148.91
saturn-noise1sp.png	1.55	165.32
saturn-noise2g.png	1.71	556.21
saturn-noise2sp.png	1.90	349.12

Table 1: Optimal values of standard deviation for which we get the most similar denoised image, in comparison to the original image; using Gaussian Filters

```

5  for i = 1.85:0.01:1.95
6      % Give noisy image and std dev as parameters to
      % gaussian filter to get a denoised image
7      filtered3 = imgaussfilt(noise1, i);
8      filtered4 = imgaussfilt(noise2, i);
9
10     % Compute errors – SSD
11     error3 = sum(sum((im - filtered3).^2));
12     error4 = sum(sum((im - filtered4).^2));
13
14     % Print results
15     fprintf('Gaussian Filter Simga %.2f– Input, Errors
        : %.2f %.2f\n', i, error3, error4)
16
17     %Display results
18     figure(2);
19     subplot(5, 8, count); imshow(filtered3); title(
        sprintf('SE %.2f', error3));
20     figure(3);
21     subplot(5, 8, count); imshow(filtered4); title(
        sprintf('SE %.2f', error4));
22     count = count + 1;
23 end

```

We report the optimal error and standard deviation σ here [See Table 1], and show the output denoised images with those optimal parameters, for better comparison with other values.

1.2 (b) Median Filtering

A median filter operates over a window by selecting the median intensity in the window.

The window of a 2D median filter can be of any central symmetric shape, a round disc, a square, a rectangle, or a cross. The pixel at the center will be replaced by the median of all pixel values inside the window.




Figure 2: A 3 x 3 median filter and how a pixel value is determined by operating around its neighborhood [1]

We use `medfilt2(I, [m n])`, which performs median filtering, where each output pixel contains the median value in the m-by-n neighborhood around the corresponding pixel in the input image.

This is shown in **line 12 and 13** of code snippet 1.2.

Matlab Implementation of Median Filters

```
1 %% Denoising algorithm (Median filtering)
2
3 % To find the std dev where there is least SSD
4 min5 = Inf; index_x5 = 0; index_y5 = 0;
5 min6 = Inf; index_x6 = 0; index_y6 = 0;
6
7 % Experiment with different filter kernel dimensions
```

```

8  for i = 1:10
9      for j = 1:10
10
11          % Convolute noisy image with i x j dimension
            median filter kernel
12          filtered5 = medfilt2(noise1, [i, j]);
13          filtered6 = medfilt2(noise2, [i, j]);
14
15          % Compute errors – SSD
16          error5 = sum(sum((im - filtered5).^2));
17          error6 = sum(sum((im - filtered6).^2));
18          fprintf('Median Filter Neighborhood %d %d–
            Input, Errors: %.2f %.2f\n', i, j, error5,
            error6)
19
20          % Save the filter dimensions of the kernel
            with the least error
21          if error5 < min5
22              min5 = error5;
23              index_x5 = i;
24              index_y5 = j;
25          end
26          if error6 < min6
27              min6 = error6;
28              index_x6 = i;
29              index_y6 = j;
30          end
31      end
32  end
33
34  % Print the results
35  fprintf('Median Filter Neighborhood %d %d– image1,
            Error: %.2f\n', index_x5, index_y5, min5)
36  fprintf('Median Filter Neighborhood %d %d– image2,
            Error: %.2f\n', index_x6, index_y6, min6)

```

We report the optimal error and filter kernel dimensions $m \times n$ here [See Table 2], and show the output denoised images with those optimal parameters, for better comparison with other values.

Image Name	Optimal Kernel Dimensions $m \times n$	Least Error - SSD
saturn-noise1g.png	5 x 7	90.11
saturn-noise1sp.png	3 x 3	8.78
saturn-noise2g.png	7 x 9	225.28
saturn-noise2sp.png	3 x 3	18.57

Table 2: Optimal values of filter kernel dimensions for which we get the most similar denoised image, in comparison to the original image; using Median Filters

1.3 (c) Non-Local Means Filtering

In this method, for each pixel in the input image the algorithm computes the denoised value as weighted mean of all the pixel values within the window, where the weight is some inverse function of the distance between patches centered at the pixels.



Figure 3: Non-Local Means Filtering - Similar pixel neighborhoods give a large weight, $w(p,q1)$ and $w(p,q2)$, while much different neighborhoods give a small weight $w(p,q3)$. [4]

Now we point out the steps of our implementation, the full matlab

code for it was too long to include here, so it can be found in Appendix A.2

- If we want a patch of size $n \times n$, its radius or half width will be $\frac{n-1}{2}$. The -1 is for the center of the patch, the pixel for which we want a $n \times n$ patch around it. So, for a 5×5 patch, the `halfPatchSize` is 2.


```

1 halfPatchSize = 2; % Patch Size (halved)
2 windowHalfSearchSize = 10; % Window Size (halved)
3 gamma = 0.01; % Bandwidth parameter

```
- Same concept as above for the window size, which is a frame made up of top, bottom, left and right edges; being certain distance away from the pixel in concern. So, for a 21×21 patch, the `windowHalfSearchSize` is 10.
- Keep in mind that, the window is NOT a patch in the sense that a patch include ALL the pixels from its left border to the right and top to the bottom. While a window is ONLY those borders.
- Each pixel of that border will have their own patches of `patchSize`, the border pixel being the center of the patch. These patches will be compared to the patch whose center is the center of the window and whose denoised value we want to compute.


```

1 % Get the patch with the pixel (i, j) as its
  center and halfPatchSize as its radius
2 px = img_n(max(1, j - halfPatchSize):min(w, j +
  halfPatchSize), max(1, i - halfPatchSize):min(h
  , i + halfPatchSize));

```
- Initially, we set bandwidth parameter γ to 0.01.
- We loop through all pixels of the image: first we get patch P_x for each pixel (x, y) by Matlab's matrix cut/crop out mechanism.
- We initialize the nominator and denominator sums to 0.0 for the denoised pixel equation

$$\hat{I}(x) = \frac{\sum_{y \in W(x)} \exp\{-\gamma \|P(x) - P(y)\|_2^2\} I(y)}{\sum_{y \in W(x)} \exp\{-\gamma \|P(x) - P(y)\|_2^2\}}$$

- Then we go through each of the four window edges around the pixel (x, y) and get the patches P_y for the pixels on the window edges, calculate the dissimilarity between P_x and P_y through SSD.
Example for Left Border:

```

1  % For each pixel on vertical edges of the window
   specified by its half size
2  for k = max(1, j - windowHalfSearchSize):min(w, j
   + windowHalfSearchSize)
3
4      % For left edge of the window
5      if 1 <= i - windowHalfSearchSize
6          % Get a patch around that
           particular pixel on the left
           edge of the window
7          py = img_n(max(1, k -
           halfPatchSize):min(w, k +
           halfPatchSize), max(1, i -
           windowHalfSearchSize -
           halfPatchSize):min(h, i -
           windowHalfSearchSize +
           halfPatchSize));
8          % Find out minimum SSD
9          if size(px) == size(py)
10             s = (px - py).^2;
11             s = s(:);
12             d = sqrt(sum(s));
13             d1 = exp(-1 * gamma * d);
14             sum_dn = sum_dn + d1;
15             sum_up = sum_up + d1 *
               img_n(j, i);
16             end
17         end
18 end

```

- We put that SSD values and other multiplication factors (bandwidth parameters and pixel values) in the given equation and after going through all four edges, get the final nominator and denominator sums.
- Dividing those will give us the denoised value for pixel (x, y) .

Image Name	Least Error - SSD
saturn-noise1g.png	102.8368
saturn-noise1sp.png	95.7399
saturn-noise2g.png	534.652
saturn-noise2sp.png	272.9644

Table 3: Optimal error values at $\gamma = 0.2$ using Non-Local Means Filters

```

1 % Get the new value for the pixel (i , j)
2 if sum_dn ~= 0
3     img_f(j , i) = sum_up/sum_dn;
4     disp(img_f(j , i))
5 end

```

We experimented with different γ values in the range of [0.01 2], we found that $\gamma = 0.2$ gives the least errors. We report the optimal error [See Table 3], and show the output denoised images with those optimal parameters, for better comparison with other values.

As we can see, non-local means filtering works better on salt and pepper noise, then it does for gaussian noise. Thorough analysis given in the next section.

1.4 Results and Conclusions

From Figures 4, 5 and 6; we can infer the following:

- **MEDIAN FILTERS** work best for salt and pepper noise, even on highly denoised image, it had an SSD between denoised and original image, of only 18.57.
- That is because salt and pepper noise only affects a pixel, the rest of the neighboring pixels do not change their values, so it is pretty easy to estimate the pixel's value, when adjacent pixels have their values intact, unlike gaussian noise, which affects all the neighbors, even if it the noise is proportional to the distance, the main takeaway is that it changed the pixel values, therefore, the whole image looks blurry as the result.
- **NON-LOCAL MEANS FILTERS** outperforms Gaussian Filters in all four cases, in some cases, from a very little margin.

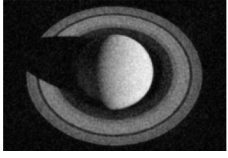
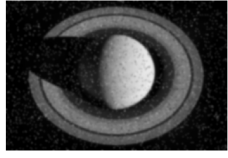
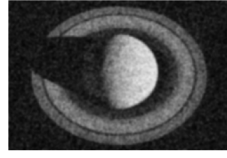
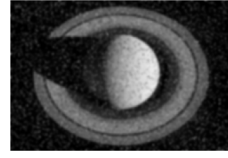
			
(a) noise1g.png denoised with gaussian filter; original noise error: 692.66, denoised error: 149.91	(b) noise1sp.png denoised with gaussian filter; original noise error: 1982.10, denoised error: 165.32	(c) noise2g.png denoised with gaussian filter; original noise error: 3118.10, denoised error: 556.21	(d) noise2sp.png denoised with gaussian filter; original noise error: 3829.28, denoised error: 349.12

Figure 4: All four noisy image denoised with gaussian filters, each having their own optimal standard deviation σ

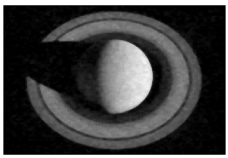
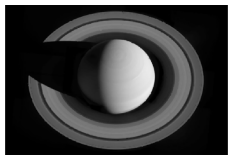
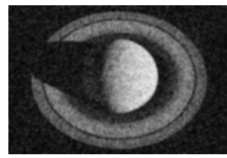
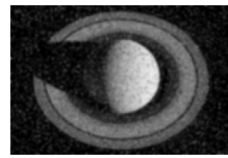
			
(a) noise1g.png denoised with median filter; original noise error: 692.66, denoised error: 90.11	(b) noise1sp.png denoised with median filter; original noise error: 1982.10, denoised error: 8.78	(c) noise2g.png denoised with median filter; original noise error: 3118.10, denoised error: 225.28	(d) noise2sp.png denoised with gaussian filter; original noise error: 3829.28, denoised error: 18.57

Figure 5: All four noisy image denoised with median filters, each having their own optimal kernel size $m \times m$

- The reason behind this can be that Non-Local Means considers weighted average of pixels around a noisy pixel(patch) with respect to other patches around its window. That gives a better global view than just considering a local neighboring pixels patch (gaussian filter).
- The reason why it is called non-local is because it does not consider the spatial relationship between pixels, instead, the similarity between different patches and the patch around the target pixel is computed to determine the weight mentioned above (inversely proportional to the distance between its patch and the reference patch of the target pixel).

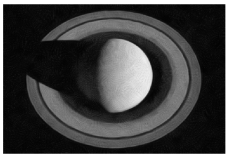
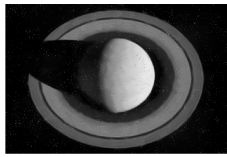
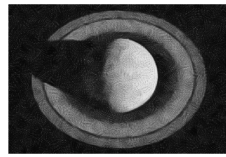
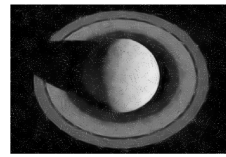
			
(a) noise1g.png denoised with Non-Local Means filter; original noise error: 692.66, denoised error: 102.8368	(b) noise1sp.png denoised with Non-Local Means filter; original noise error: 1982.10, denoised error: 95.7399	(c) noise2g.png denoised with Non-Local Means filter; original noise error: 3118.10, denoised error: 534.652	(d) noise2sp.png denoised with Non-Local Means filter; original noise error: 3829.28, denoised error: 272.9644

Figure 6: All four noisy image denoised with Non-Local Means filters, each having their bandwidth parameter $\gamma = 0.2$

- **MEDIAN FILTERS** gives better results for gaussian noised images too.
- Median filter (rectangular kernel) is optimal for reducing random noise in spatial domain (image space). However Median filter is the worst filter for frequency domain, with little ability to separate one band of frequencies from another. Gaussian filter has better performance in frequency domain.
- Gaussian filter uses convolution and is very slow. If we implement Mean filter using recursive formula it will run like very fast.
- **MEAN FILTER IS FAST AND PROBABLY THE BEST SOLUTION IF WE WANT TO REMOVE NOISE FROM IMAGE. IT IS BAD SOLUTION IF YOU WANT TO SEPARATE FREQUENCIES PRESENT IN THE IMAGE.**

2 Problem 2: Texture Synthesis

TEXTURE SYNTHESIS is the process of algorithmically constructing a large digital image from a small digital sample image by taking advantage of its structural content.

The patterns to be synthesized can be **REGULAR** - repeating in fixed spaces in a fixed/ replicated way; the patterns can be **STOCHASTIC** - they look like noise; or they can be **IRREGULAR** - they have no specific recurring structure at all.

2.1 Objective

Our goal for this mini-project is to[5]:

- The output should have the **SIZE** given by the user.
- The output should be as **SIMILAR** as possible to the sample.
- The output should **NOT HAVE VISIBLE ARTIFACTS** such as seams, blocks and misfitting edges.
- The output should not **REPEAT**, i. e. the same structures in the output image should not appear multiple places.
- The process of texture synthesis should be efficient in **COMPUTATION TIME AND IN MEMORY USE**.

Now, we look at two methods to achieve our goal, in the next sections.

2.2 (a) Random Tiling

The procedure was simple for random tiling; as there is no logic except to pick up the random tiles from within the input image bounds and to put them on the correct position.

The Matlab code for this is in Appendix B.1 [for the starting point of the program, which calls `synthRandomPatch.m`] and B.2 [for the function `synthRandomPatch` implementation].

Again, we took care of not choosing a random co-ordinate pair in a way that might lead to overrunning the input image bounds and thus, not having the exact same size as the pre-decided `tileSize`.

If that was the case, we simply chose some other random co-ordinate pair which was within the bounds.

2.2.1 Results

Figures 7, 8 and 9 shows the result of random sampling for four tile sizes - 15, 20, 30, 40.

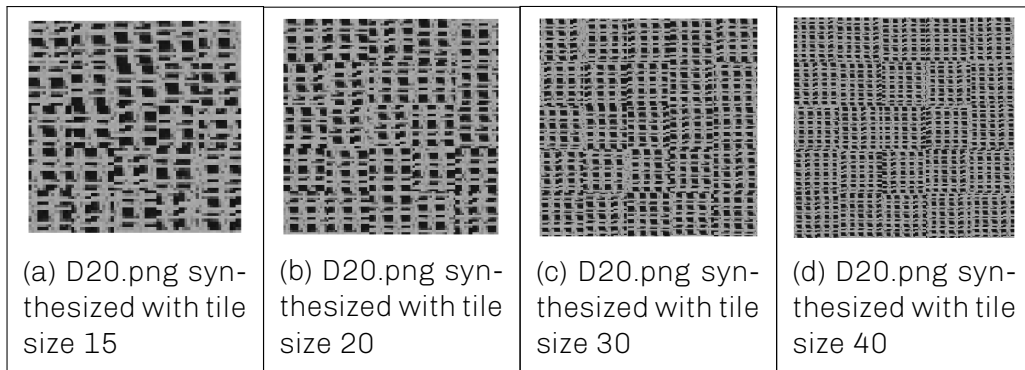


Figure 7: D20.png synthesized with tile sizes - 15, 20, 30, 40. Note that the original sizes of images are 75x75, 100x100, 150x150, 200x200; but we are showing them on a same scale for better representation and understanding of how they look side by side

As we can see, only in an extremely rare case, this method will work, when the chosen tiles are always seamless. Even in a fixed pattern, choosing a tile randomly in this way is near to impossible.

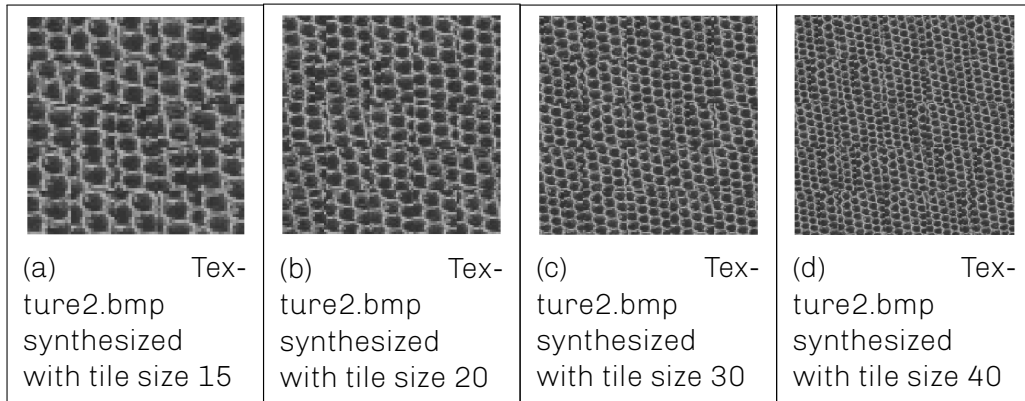


Figure 8: Texture2.bmp synthesized with tile sizes - 15, 20, 30, 40. Note that the original sizes of images are 75x75, 100x100, 150x150, 200x200; but we are showing them on a same scale for better representation and understanding of how they look side by side

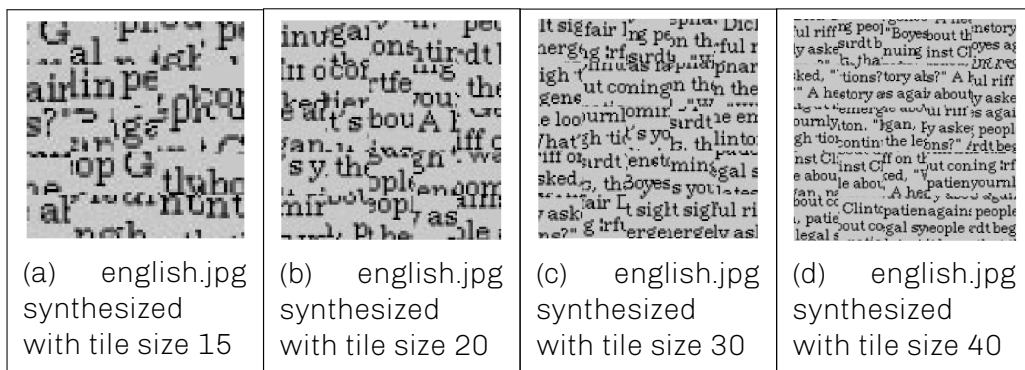


Figure 9: english.jpg synthesized with tile sizes - 15, 20, 30, 40. Note that the original sizes of images are 75x75, 100x100, 150x150, 200x200; but we are showing them on a same scale for better representation and understanding of how they look side by side

Therefore, we certainly need a better method which promises a smooth transition from one to the next tile. But random tiling provides us with the basic knowledge to approach the texture synthesis problem.

2.3 (b) Non-parametric Sampling

In Non-parametric Sampling, the texture synthesis process grows a new image outward from **AN INITIAL SEED**, one pixel at a time.

A Markov random field model is assumed, and the conditional distribution of a pixel given all its neighbors synthesized so far is estimated by querying the sample image and **FINDING ALL SIMILAR NEIGHBORHOODS**.

The method aims at **PRESERVING AS MUCH LOCAL STRUCTURE** as possible and produces good results for a wide variety of synthetic and real-world textures.

We want to insert pixel intensities **BASED ON EXISTING NEARBY PIXEL** values.

Distribution of a value of a pixel is conditioned on its neighbors alone.

The Matlab code for this is in Appendix B.1 [for the starting point of the program, which calls `synthEfrosLeung.m`] and B.3 [for the function `synthEfrosLeung` implementation].

2.3.1 Outputs

We ran our algorithm 4-5 times for each window size and for each image, so we can average the time and decrease the effect of “worse” choice of seed, even if it is just by a very little margin.

We display the best result [more realistic and reasonable result] out of all the trials.

Figures 10, 11 and 12 shows the result of random sampling for four tile sizes - 5, 7, 11, 15.

2.3.2 Effect of window size on runtime and on synthesis quality

We note down our observation about correlation between window sizes and runtime in Table 4:

We now note down few of our inferences about correlation between window sizes, synthesis quality and runtime:

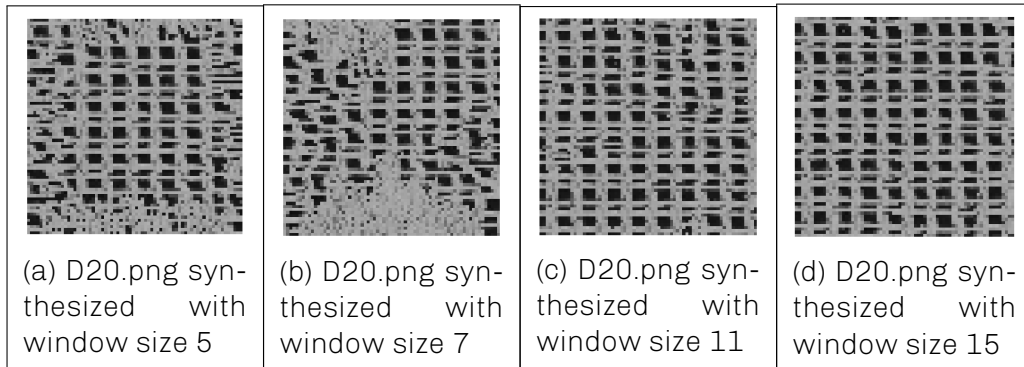


Figure 10: D20.png synthesized with window sizes - 5, 7, 11, 15. Note that the sizes of all output images are 70x70

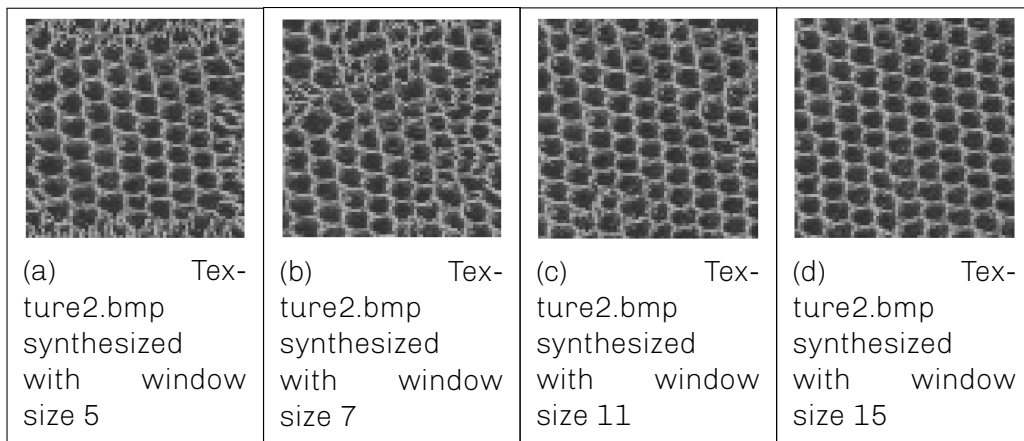


Figure 11: Texture2.bmp synthesized with window sizes - 5, 7, 11, 15. Note that the sizes of all output images are 70x70

- We can observe from the Table 4 that as the size of the window increases, the runtime increases as the patch sizes will be bigger and thus, the distance metric computations also takes longer.
- But, since the bigger window size means more neighborhood coverage, the windows sizes of 11 and 15 gives a lot more realistic results then the windows of 5 and 7.
- That is because when we look at a small patch, there might be patterns which match with that tiny patch but globally, they might have no correlation at all.
- That makes a particular pixel with tiny patch get “wrong” value, which

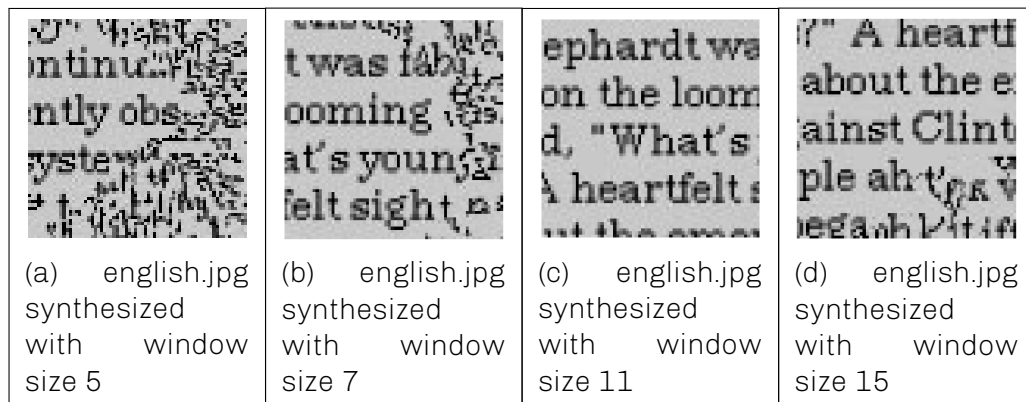


Figure 12: english.jpg synthesized with window sizes - 5, 7, 11, 15. Note that the sizes of all output images are 70x70

just snowballs upon further synthesis.

- **IN CONCLUSION** , larger window sizes will give better quality of texture synthesis since they have better "global coverage", but it also takes more runtime for the same reason [more computation for distance metric because more pixels in the patch]

Image Name	Window Size	Average Running Time for 5 Runs (in seconds)
D20.png Input Size: 49 x 53 Output Size: 70 x 70	5	18.762
	7	28.449
	11	49.166
	15	73.546
Texture2.bmp Input Size: 58 x 56 Output Size: 70 x 70	5	19.569
	7	34.032
	11	67.694
	15	81.865
english.jpg Input Size: 151 x 156 Output Size: 70 x 70	5	137.006
	7	207.111
	11	489.050
	15	786.656

Table 4: Relation between window size, input image size and the run-time

3 Extensions for Extra Credit

3.1 Extension 1: Block Matching 3D (BM3D)

BLOCK-MATCHING AND 3D FILTERING (BM3D) is a 3-D block-matching algorithm used primarily for noise reduction in images.[6]

- This strategy is based on an enhanced **SPARSE REPRESENTATION** in transform-domain. The enhancement of the sparsity is achieved by grouping similar 2D image fragments (e.g. blocks) into 3D data arrays which we call "groups".
- Collaborative filtering is a special procedure developed to deal with these 3D groups. It is realized by using the three successive steps: 3D **TRANSFORMATION** of 3D group, **SHRINKAGE** of transform spectrum, and **INVERSE** 3D transformation.
- The result is a 3D estimate that consists of the **JOINTLY FILTERED** grouped image blocks.
- By attenuating the noise, the collaborative filtering reveals even the finest details shared by grouped blocks and at the same time it preserves the essential **UNIQUE FEATURES** of each individual block.
- The filtered blocks are then returned to their original positions.
- Because these **BLOCKS ARE OVERLAPPING**, for each pixel we obtain many different estimates which need to be combined.
- **AGGREGATION** is a particular averaging procedure which is exploited to take advantage of this redundancy.
- A significant improvement is obtained by a specially developed collaborative **WIENER FILTERING**.
- The results presented here demonstrate that the developed methods achieve state-of-the-art denoising performance in terms of both peak signal-to-noise ratio and subjective visual quality.

The PSNR results are shown in 5 and the output images are shown in Figures 13.

We now compare BM3D with NLM from Section 1.3:

Image Name	Optimal Standard Deviation σ	Peak Signal-to-Noise Ratio PSNR	SSD
saturn-noise1g.png	0.5	21.418	203.346
saturn-noise1sp.png	0.3	16.852	370.128
saturn-noise2g.png	0.1	14.8844	482.735
saturn-noise2sp.png	0.1	13.9922	568.804

Table 5: Optimal values of standard deviation for which we get the most similar denoised image, in comparison to the original image; using Block-Matching 3D Filtering

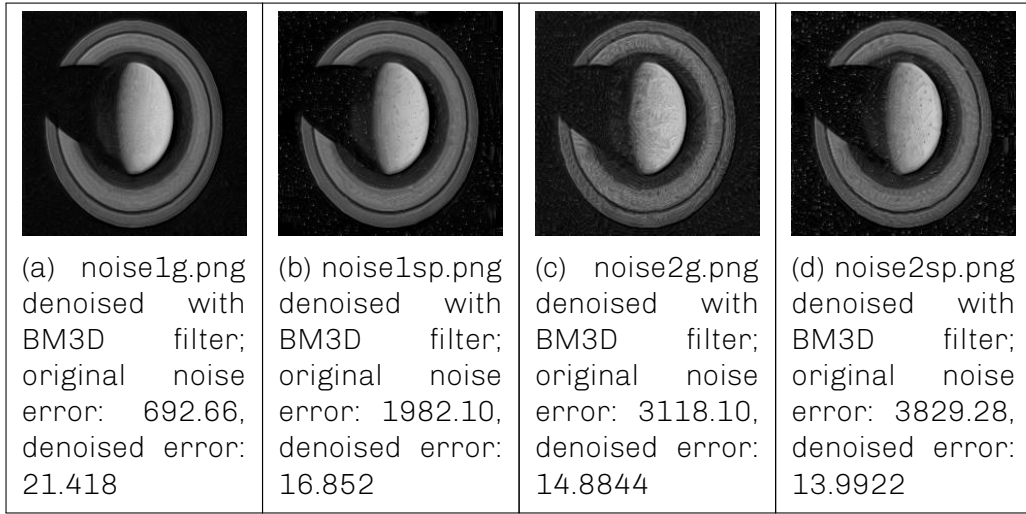


Figure 13: All four noisy image denoised with BM3D filters, each having their own optimal standard deviation σ

- Our NLM representation still outperforms BM3D implementation of Kostadin Dabov [6], see Tables 3 and 5
- BM3D takes the concept of non-local means (NLM) in the sense that it also attempts denoising based on finding similar patches within a given window.
- BM3D use a reduced set of adjacent pixels, but in comparison to NLM, it also requires additional, computationally intensive processing steps and continue to suffer from visible artifacts on sharp edges and in smooth regions.
- Hence, the sub-optimal BM3D results as opposed to quite satisfactory results from NLM.

- BM3D approach generates a slightly more noisy image, but with better preserved finer details. And without introducing grid artifacts that are on the other hand visible.

BM3D adopts a NLM search selection strategy, but it also suffers from the noise-to-noise matching problem. That's why we get better results with 3 out 4 noisy images, with BM3D performing only better at large gaussian noise.

Although, there exists better BM3D implementation then the one we used here; which outperforms NLM by a large margin. E.g., BM4D, VBM3D, VBM4D, RF3D, BM3D-CFA, BM3D-SAPCA

3.2 Extension 2: Image Quilting

IMAGE QUILTING proposes a way to minimize these by picking tiles that align well along the boundary (like solving a jigsaw puzzle). The basic idea is to pick a tile that has small SSD along the overlapping boundary with the tiles generated so far.

Here we show the results of Image Quilting techniques described by Efros and Freeman.

We select the first seed random tile as shown below:

```
1 %% Get random tile co-ordinate pairs for the seed
   patch
2 x = floor(1 + rand(1,1) * (w - tileSize - 1));
3 y = floor(1 + rand(1,1) * (h - tileSize - 1));
4
5 % Get the seed patch
6 patch = im(x:x+tileSize - 1, y:y+tileSize - 1);
7 % Copy the random tile seed patch from source image to
   the top-left corner
8 im_patch(1:tileSize , 1:tileSize) = patch;
```

We select the overlapping factor as shown below, if it is floating point, we round it off to a smallest integer greater than the float:

```
1 %% size of the region of overlap between tiles –
   typically 1/6
2 overlap = ceil(1/6 * tileSize);
```

We select the overlapping region for vertical and horizontal parts as shown below:

```
1 % Upper Horizontal overlapping section
2 A = temp_patch(1:tileSize , 1:overlap);
3 % Left Vertical overlapping section
4 B = temp_patch(1:overlap , 1:tileSize );
```

The **MATLAB CODE** for this is in Appendix B.1 [for the starting point of

the program, which calls `synthEfrosLeung.m`] and C.1 [for the function `synthImageQuilting` implementation].

The results are shown in Figures 14, 15 and 16.

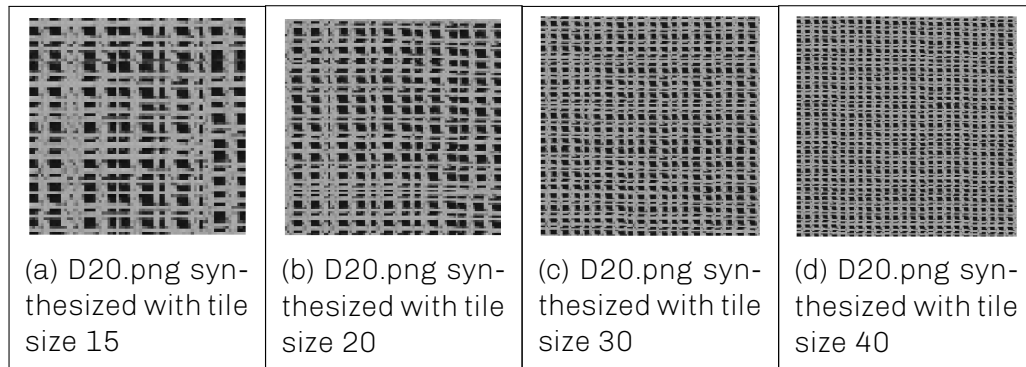


Figure 14: D20.png quilted with tile sizes - 15, 20, 30, 40. Note that the original sizes of images are 75x75, 100x100, 150x150, 200x200; but we are showing them on a same scale for better representation and understanding of how they look side by side

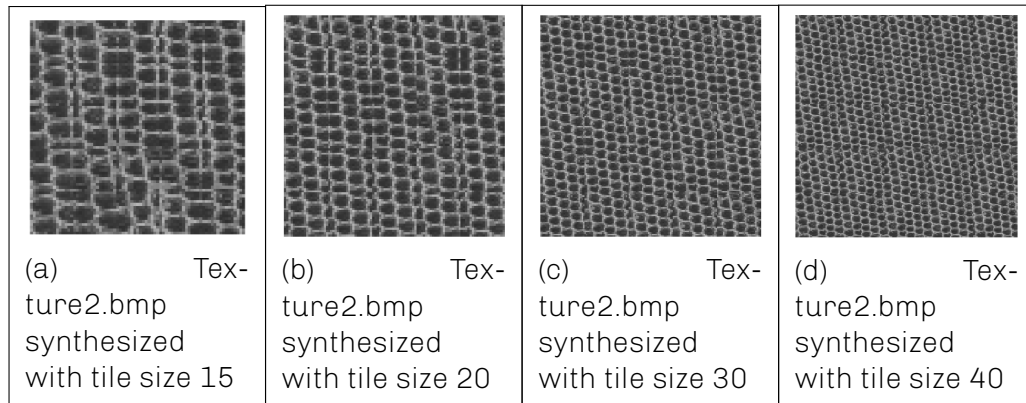


Figure 15: Texture2.bmp quilted with tile sizes - 15, 20, 30, 40. Note that the original sizes of images are 75x75, 100x100, 150x150, 200x200; but we are showing them on a same scale for better representation and understanding of how they look side by side

As we can see, the text image `english.jpg` still needs to be adjusted to proper window size, in order to address the doubling/repeating effect.

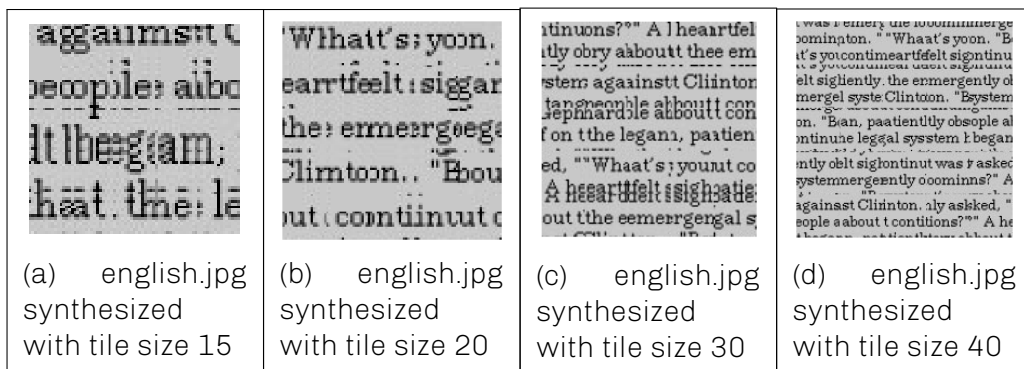


Figure 16: english.jpg quilted with tile sizes - 15, 20, 30, 40. Note that the original sizes of images are 75x75, 100x100, 150x150, 200x200; but we are showing them on a same scale for better representation and understanding of how they look side by side

Notwithstanding its plainness, this technique works astoundingly good when tried for texture synthesis, fabricating results that are equal or better than the Efros & Leung family of algorithms but with upgraded strength (less chance of “growing garbage”) and at a snippet of the computational price.

References

- [1] Subhransu Maji, *Linear Filtering* https://www.dropbox.com/s/tfsmuiunffyu7ml/lec09\%2B10_linear_filtering.pdf?dl=0
- [2] *Noise Reduction* https://en.wikipedia.org/wiki/Noise_reduction#In_images
- [3] Richard Szeliski, *Computer Vision: Algorithms and Applications* ISBN: 978-1-848-82934-3
- [4] Antoni Buades, Bartomeu Coll, Jean-Michel Morel, *A non-local algorithm for image denoising* <https://www.iro.umontreal.ca/~mignotte/IFT6150/Articles/Buades-NonLocal.pdf>
- [5] *Texture synthesis* https://en.wikipedia.org/wiki/Texture_synthesis
- [6] *Block-matching and 3D filtering (BM3D) algorithm and its extensions* http://www.cs.tut.fi/~foi/GCF-BM3D/index.html#ref_problems

A

Matlab code for Problem 1: Image Denoising

A.1 Implementation of evalDenoising.m

```
1 % Entry code for evaluating demosaicing algorithms
2 % The code loops over all images and methods, computes
  the error and
3 % displays them in a table.
4 %
5 %
6 % This code is part of:
7 %
8 %   CMPSCI 670: Computer Vision
9 %   University of Massachusetts, Amherst
10 %   Instructor: Subhransu Maji
11 %
12 close all;
13 clc;
14 % Load images
15 im = im2double(imread(' ../data/denoising/saturn.png'))
    ;
16 noise1 = im2double(imread(' ../data/denoising/saturn-
    noise1g.png'));
17 noise2 = im2double(imread(' ../data/denoising/saturn-
    noise1sp.png'));
18 %noise1 = im2double(imread(' ../data/denoising/saturn-
    noise2g.png'));
19 %noise2 = im2double(imread(' ../data/denoising/saturn-
    noise2sp.png'));
20
21 % Compute errors
22 error1 = sum(sum((im - noise1).^2));
23 error2 = sum(sum((im - noise2).^2));
24 fprintf('Input, Errors: %.2f %.2f\n', error1, error2)
25
26 % Display the images
27 figure(1);
```

```

28 subplot(1,3,1); imshow(im); title('Input');
29 subplot(1,3,2); imshow(noise1); title(sprintf('SE %.2f
    ', error1));
30 subplot(1,3,3); imshow(noise2); title(sprintf('SE %.2f
    ', error2));
31
32 %% Denoising algorithm (Gaussian filtering)
33
34 count = 1;
35 % Experiment with different values of standard
    deviation
36 for i = 1.85:0.01:1.95
37     % Give noisy image and std dev as parameters to
        gaussian filter to get a denoised image
38     filtered3 = imgaussfilt(noise1, i);
39     filtered4 = imgaussfilt(noise2, i);
40
41     % Compute errors – SSD
42     error3 = sum(sum((im - filtered3).^2));
43     error4 = sum(sum((im - filtered4).^2));
44
45     % Print results
46     fprintf('Gaussian Filter Simga %.2f– Input, Errors
        : %.2f %.2f\n', i, error3, error4)
47
48     %Display results
49     figure(2);
50     subplot(5, 8, count); imshow(filtered3); title(
        sprintf('SE %.2f', error3));
51     figure(3);
52     subplot(5, 8, count); imshow(filtered4); title(
        sprintf('SE %.2f', error4));
53     count = count + 1;
54
55
56 %% Denoising algorithm (Median filtering)
57
58 % To find the std dev where there is least SSD
59 min5 = Inf; index_x5 = 0; index_y5 = 0;
60 min6 = Inf; index_x6 = 0; index_y6 = 0;
61

```

```

62 % Experiment with different filter kernel dimensions
63 for i = 1:10
64     for j = 1:10
65
66         % Convolute noisy image with i x j dimension
        % median filter kernel
67         filtered5 = medfilt2(noise1, [i, j]);
68         filtered6 = medfilt2(noise2, [i, j]);
69
70         % Compute errors – SSD
71         error5 = sum(sum((im - filtered5).^2));
72         error6 = sum(sum((im - filtered6).^2));
73         fprintf('Median Filter Neighborhood %d %d–
            Input, Errors: %.2f %.2f\n', i, j, error5,
            error6)
74
75         % Save the filter dimensions of the kernel
        % with the least error
76         if error5 < min5
77             min5 = error5;
78             index_x5 = i;
79             index_y5 = j;
80         end
81         if error6 < min6
82             min6 = error6;
83             index_x6 = i;
84             index_y6 = j;
85         end
86     end
87 end
88
89 % Print the results
90 fprintf('Median Filter Neighborhood %d %d– image1,
        Error: %.2f\n', index_x5, index_y5, min5)
91 fprintf('Median Filter Neighborhood %d %d– image2,
        Error: %.2f\n', index_x6, index_y6, min6)
92
93 %% Denoising algorithm (Non-local means)
94
95 img_n = noise1;      % Noisy Image
96 halfPatchSize = 2;   % Patch Size (halved)

```

```

97 windowHalfSearchSize = 10; % Window Size (halved)
98 gamma = 0.01; % Bandwidth parameter
99 [w, h] = size(img_n); % Size of the noisy image
100
101 img_f = img_n; % Denoised Image
102
103 % Go through whole image
104 for i = 1:h
105     for j = 1:w
106
107         % Get the patch with the pixel (i, j) as its
            center and halfPatchSize as its radius
108         px = img_n(max(1, j - halfPatchSize):min(w, j
            + halfPatchSize), max(1, i - halfPatchSize)
            :min(h, i + halfPatchSize));
109
110         % nominator of the denoised pixel equation
111         sum_up = 0.0;
112         % denominator of the denoised pixel equation
113         sum_dn = 0.0;
114
115         % For each pixel on vertical edges of the
            window specified by its half size
116         for k = max(1, j - windowHalfSearchSize):min(w
            , j + windowHalfSearchSize)
117
118             % For left edge of the window
119             if 1 <= i - windowHalfSearchSize
120                 % Get a patch around that particular
                    pixel on the left edge of the
                    window
121                 py = img_n(max(1, k - halfPatchSize):
                    min(w, k + halfPatchSize), max(1, i
                    - windowHalfSearchSize -
                    halfPatchSize):min(h, i -
                    windowHalfSearchSize +
                    halfPatchSize));
122                 % Find out minimum SSD
123                 if size(px) == size(py)
124                     s = (px - py).^2;
125                     s = s(:);

```

```

126         d = sqrt(sum(s));
127         d1 = exp(-1 * gamma * d);
128         sum_dn = sum_dn + d1;
129         sum_up = sum_up + d1 * img_n(j, i)
            ;
130     end
131 end
132
133 % For right edge of the window
134 if h >= i + windowHalfSearchSize
135     % Get a patch around that particular
        pixel on the right edge of the
        window
136     py = img_n(max(1, k - halfPatchSize):
        min(w, k + halfPatchSize), max(1, i
        + windowHalfSearchSize -
        halfPatchSize):min(h, i +
        windowHalfSearchSize +
        halfPatchSize));
137     % Find out minimum SSD
138     if size(px) == size(py)
139         s = (px - py).^2;
140         s = s(:);
141         d = sqrt(sum(s));
142         d1 = exp(-1 * gamma * d);
143         sum_dn = sum_dn + d1;
144         sum_up = sum_up + d1 * img_n(j, i)
            ;
145     end
146 end
147 end
148
149 % For each pixel on horizontal edges of the
        window specified by its half size
150 for k = max(1, i - windowHalfSearchSize):min(h
        , i + windowHalfSearchSize)
151     % For top edge of the window
152     if 1 <= j - windowHalfSearchSize
153         % Get a patch around that particular
            pixel on the upper edge of the
            window

```



```

154         py = img_n(max(1, j -
            windowHalfSearchSize -
            halfPatchSize):min(w, j -
            windowHalfSearchSize +
            halfPatchSize), max(1, k -
            halfPatchSize):min(h, k +
            halfPatchSize));
155     % Find out minimum SSD
156     if size(px) == size(py)
157         s = (px - py).^2;
158         s = s(:);
159         d = sqrt(sum(s));
160         d1 = exp(-1 * gamma * d);
161         sum_dn = sum_dn + d1;
162         sum_up = sum_up + d1 * img_n(j, i)
            ;
163     end
164 end
165
166 % For bottom edge of the window
167 if w >= j + windowHalfSearchSize
168     % Get a patch around that particular
        pixel on the lower edge of the
        window
169     py = img_n(max(1, j +
        windowHalfSearchSize -
        halfPatchSize):min(w, j +
        windowHalfSearchSize +
        halfPatchSize), max(1, k -
        halfPatchSize):min(h, k +
        halfPatchSize));
170     % Find out minimum SSD
171     if size(px) == size(py)
172         s = (px - py).^2;
173         s = s(:);
174         d = sqrt(sum(s));
175         d1 = exp(-1 * gamma * d);
176         sum_dn = sum_dn + d1;
177         sum_up = sum_up + d1 * img_n(j, i)
            ;
178     end

```

```

179         end
180     end
181
182     % Get the new value for the pixel (i, j)
183     if sum_dn ~= 0
184         img_f(j, i) = sum_up/sum_dn;
185         disp(img_f(j, i))
186     end
187
188 end
189 end
190
191 % Caluclate Error – SSD
192 error7 = sum(sum((im - img_f).^2));
193 % Print the error
194 fprintf('Non-Linear Means Errors: %.6f\n', error7)
195 % Show the results
196 figure;
197 subplot(121); imshow(im);
198 subplot(122); imshow(img_f);

```

A.2 Implementation of Non-Local Means Filters

NOTE: this is included in the `evalDenoising.m` shown in A.1 too, the separate copy is just for clarification.

```

1 %% Denoising alogirthm (Non-local means)
2
3 img_n = noise1;      % Noisy Image
4 halfPatchSize = 2;  % Patch Size (halved)
5 windowHalfSearchSize = 10; % Window Size (halved)
6 gamma = 0.01;       % Bandwidth parameter
7 [w, h] = size(img_n); % Size of the noisy image
8
9 img_f = img_n;      % Denoised Image
10
11 % Go through whole image
12 for i = 1:h
13     for j = 1:w
14
15         % Get the patch with the pixel (i, j) as its
           center and halfPatchSize as its radius

```

```

16     px = img_n(max(1, j - halfPatchSize):min(w, j
        + halfPatchSize), max(1, i - halfPatchSize)
        :min(h, i + halfPatchSize));
17
18     % nominator of the denoised pixel equation
19     sum_up = 0.0;
20     % denominator of the denoised pixel equation
21     sum_dn = 0.0;
22
23     % For each pixel on vertical edges of the
        window specified by its half size
24     for k = max(1, j - windowHalfSearchSize):min(w
        , j + windowHalfSearchSize)
25
26         % For left edge of the window
27         if 1 <= i - windowHalfSearchSize
28             % Get a patch around that particular
                pixel on the left edge of the
                window
29             py = img_n(max(1, k - halfPatchSize):
                min(w, k + halfPatchSize), max(1, i
                - windowHalfSearchSize -
                halfPatchSize):min(h, i -
                windowHalfSearchSize +
                halfPatchSize));
30             % Find out minimum SSD
31             if size(px) == size(py)
32                 s = (px - py).^2;
33                 s = s(:);
34                 d = sqrt(sum(s));
35                 d1 = exp(-1 * gamma * d);
36                 sum_dn = sum_dn + d1;
37                 sum_up = sum_up + d1 * img_n(j, i)
                    ;
38             end
39         end
40
41         % For right edge of the window
42         if h >= i + windowHalfSearchSize
43             % Get a patch around that particular
                pixel on the right edge of the

```

```

44         window
py = img_n(max(1, k - halfPatchSize):
min(w, k + halfPatchSize), max(1, i
+ windowHalfSearchSize -
halfPatchSize):min(h, i +
windowHalfSearchSize +
halfPatchSize));
45 % Find out minimum SSD
46 if size(px) == size(py)
47     s = (px - py).^2;
48     s = s(:);
49     d = sqrt(sum(s));
50     d1 = exp(-1 * gamma * d);
51     sum_dn = sum_dn + d1;
52     sum_up = sum_up + d1 * img_n(j, i)
        ;
53     end
54 end
55 end
56
57 % For each pixel on horizontal edges of the
    window specified by its half size
58 for k = max(1, i - windowHalfSearchSize):min(h
, i + windowHalfSearchSize)
59     % For top edge of the window
60     if 1 <= j - windowHalfSearchSize
61         % Get a patch around that particular
            pixel on the upper edge of the
            window
62         py = img_n(max(1, j -
windowHalfSearchSize -
halfPatchSize):min(w, j -
windowHalfSearchSize +
halfPatchSize), max(1, k -
halfPatchSize):min(h, k +
halfPatchSize));
63         % Find out minimum SSD
64         if size(px) == size(py)
65             s = (px - py).^2;
66             s = s(:);
67             d = sqrt(sum(s));

```

```

68         d1 = exp(-1 * gamma * d);
69         sum_dn = sum_dn + d1;
70         sum_up = sum_up + d1 * img_n(j, i)
71         ;
72     end
73
74     % For bottom edge of the window
75     if w >= j + windowHalfSearchSize
76         % Get a patch around that particular
77         % pixel on the lower edge of the
78         % window
79         py = img_n(max(1, j +
80             windowHalfSearchSize -
81             halfPatchSize):min(w, j +
82             windowHalfSearchSize +
83             halfPatchSize), max(1, k -
84             halfPatchSize):min(h, k +
85             halfPatchSize));
86         % Find out minimum SSD
87         if size(px) == size(py)
88             s = (px - py).^2;
89             s = s(:);
90             d = sqrt(sum(s));
91             d1 = exp(-1 * gamma * d);
92             sum_dn = sum_dn + d1;
93             sum_up = sum_up + d1 * img_n(j, i)
94             ;
95         end
96     end
97 end
98
99 % Get the new value for the pixel (i, j)
100 if sum_dn ~= 0
101     img_f(j, i) = sum_up/sum_dn;
102     disp(img_f(j, i))
103 end
104
105 end
106
107 end
108
109 end

```

```

99 % Caluclate Error – SSD
100 error7 = sum(sum((im - img_f).^2));
101 % Print the error
102 fprintf('Non-Linear Means Errors: %.6f\n', error7)
103 % Show the results
104 figure;
105 subplot(121); imshow(im);
106 subplot(122); imshow(img_f);

```

B

Matlab code for Problem 2: Texture Synthesis

B.1 Implementation of evalTextureSynth.m

```
1 close ;
2
3 % Load images
4 % im = im2double(imread(' ../ data/texture/D20.png' ));
5 % im = im2double(imread(' ../ data/texture/Texture2.bmp
   ' ));
6 im = im2double(imread(' ../ data/texture/english.jpg' ))
   ;
7
8 figure(1); imshow(im);
9
10
11 %% Random patches
12 tileSize = 40; % specify block sizes – 15 20 30 40
13 numTiles = 5;
14 outSize = numTiles * tileSize; % calculate output
   image size
15 % save the random-patch output and record run-times
16 im_patch = synthRandomPatch(im, tileSize, numTiles,
   outSize);
17 figure;
18 imshow(im_patch);
19
20 %% Non-parametric Texture Synthesis using Efros &
   Leung algorithm
21 winsize = 11; % specify window size (5, 7, 11, 15)
22 outSize = 70; % specify size of the output image to
   be synthesized (square for simplicity)
23 % save the synthesized image and record the run-times
24 im_synth = synthEfrosLeung(im, winsize, outSize);
25 figure(3);
26 imshow(im_synth);
27
```

```

28 %% Image Quilting
29 tileSize = 40; % specify block sizes – 15 20 30 40
30 numTiles = 5;
31 outSize = numTiles * tileSize; % calculate output
    image size
32 % save the random-patch output and record run-times
33 im_patch = synthImageQuilting(im, tileSize, numTiles,
    outSize);
34 figure;
35 imshow(im_patch);

```

B.2 Implementation of synthRandomPatch.m

```

1 function im_patch = synthRandomPatch(im, tileSize,
    numTiles, outSize)
2
3     % Get size of original image whose texture we want
    to synthesize
4     [w, h] = size(im);
5
6     % Initialize output/ result to the size specified
7     im_patch = zeros(outSize);
8
9     % Create a synthesized image block-by-block or
    tile-by-tile
10    for i = 1:numTiles
11        for j = 1:numTiles
12
13            % Randomly selects an x-axis as a top-left
                corner of a random tile
14            x = floor(1 + rand(1,1) * (w - tileSize
                -1));
15            % Randomly selects an y-axis as a top-left
                corner of a random tile
16            y = floor(1 + rand(1,1) * (h - tileSize
                -1));
17
18            % Cut out the copy of the tile
19            patch = im(x:x+tileSize-1, y:y+tileSize-1)
                ;
20            % Paste it in the assigned top-left corner

```



```

21         of the output image
22         % [Top to Bottom; Left to Right]
23         im_patch((i-1)*tileSize+1:(i-1)*tileSize+
24                 tileSize , (j-1)*tileSize+1:(j-1)*
25                 tileSize+tileSize) = patch;
26     end
27 end
28 end

```

B.3 Implementation of synthEfrosLeung.m

```

1 function im_synth = synthEfrosLeung(im, winsize,
   outSize)
2     %% Initializing output/synthesized image with all
   -1 (not with 0s as they might depict an actual
   value)
3     im_synth = zeros(outSize)-1;
4     % Maximum number of pixels which might need to get
   their values interpolated
5     n = outSize * outSize;
6
7     % Get half window size
8     halfWinSize = floor(winsize/2);
9     % Error Threshold for finding best matches to
   interpolate a pixel's value
10    error_threshold = 0.1;
11
12    %% Get input image size
13    [w, h, c] = size(im);
14
15    % Randomly choose a pixel value co-ordinate pair;
   "-3" shows we must choose the pixel in a way
   that a 3 x 3 patch with this pixel at top-left
   is possible
16    x = floor(1 + rand(1,1) * (w - 3 - 1));
17    y = floor(1 + rand(1,1) * (h - 3 - 1));
18    % A 3x3 SEED PATCH
19    patch = im(x:x+3-1, y:y+3-1);
20
21    % Copy the seed patch to the center of the output/
   synthesized image

```

```

22     im_synth(outSize/2-1:outSize/2+1, outSize/2-1:
        outSize/2+1) = patch;
23
24     %% Growing the borders of the filled pixels in the
        output image
25     for k = 1:outSize*outSize
26         %% list of pixel positions in the output image
            that are unfilled, but contain filled
            pixels in their neighbourhood
27         pixelList = zeros(n, 3);
28         count = 1;
29
30         % For all pixels of the output/synthesized
            image
31         for i = 1:outSize
32             for j = 1:outSize
33                 flag = 0;
34
35                 % For any pixel whose value has yet to
                    be calculated
36                 if im_synth(i, j) == -1
37                     % Add it to the list but...
38                     pixelList(count, 1) = i;
39                     pixelList(count, 2) = j;
40
41                     % Identify how many of its 8
                        neighbors have their values
                        interpolated already
42                     % Add the #
                        neighbors_having_valid_values to
                        the pixelList frequency counter
43                     if i + 1 <= outSize && im_synth(i
                        +1, j) ~= -1
44                         flag = 1;
45                         pixelList(count, 3) = pixelList
                            (count, 3) + 1;
46                     end
47                     if i - 1 >= 1 && im_synth(i-1, j)
                        ~= -1
48                         flag = 1;
49                         pixelList(count, 3) = pixelList

```

```

                    (count, 3) + 1;
50     end
51     if j + 1 <= outSize && im_synth(i,
        j+1) ~= -1
52         flag = 1;
53         pixelList(count, 3) = pixelList
            (count, 3) + 1;
54     end
55     if j - 1 >= 1 && im_synth(i, j-1)
        ~= -1
56         flag = 1;
57         pixelList(count, 3) = pixelList
            (count, 3) + 1;
58     end
59     if i + 1 <= outSize && j - 1 >= 1
        && im_synth(i+1, j-1) ~= -1
60         flag = 1;
61         pixelList(count, 3) = pixelList
            (count, 3) + 1;
62     end
63     if i - 1 >= 1 && j - 1 >= 1 &&
        im_synth(i-1, j-1) ~= -1
64         flag = 1;
65         pixelList(count, 3) = pixelList
            (count, 3) + 1;
66     end
67     if i + 1 <= outSize && j + 1 <=
        outSize && im_synth(i+1, j+1) ~=
        -1
68         flag = 1;
69         pixelList(count, 3) = pixelList
            (count, 3) + 1;
70     end
71     if i - 1 >= 1 && j + 1 <= outSize
        && im_synth(i-1, j+1) ~= -1
72         flag = 1;
73         pixelList(count, 3) = pixelList
            (count, 3) + 1;
74     end
75
76     % If the output pixel has at least

```

```

                                one neighbor having valid value,
                                only then that pixel would
                                count
77         if flag == 1
78             count = count + 1;
79         end
80     end
81 end
82 end
83
84 %% Now we have a list of all the pixels whose
    values we will interpolate in this pass
85 pixelList = pixelList(1:count-1,:);
86 % Sort the list by the number of filled
    neighborhood pixels
87 pixelList = sortrows(pixelList, 3, 'descend');
88
89 %% For each pixel in pixelList
90 for p = 1:size(pixelList, 1)
91     % Initialize SSD to maximum integer
92     ssd_min = Inf;
93
94     %% Initialize a validMask of windowSize x
        windowSize
95     validMask = im_synth(max(1, pixelList(p,
        1) - halfWinSize): min(outSize,
        pixelList(p, 1) + halfWinSize), max(1,
        pixelList(p, 2) - halfWinSize): min(
        outSize, pixelList(p, 2) + halfWinSize)
        );
96     % 1s at the neighborhood positions of
        that pixel that contains filled pixels
97     validMask(validMask ~= -1) = 1;
98     % 0s for unfilled positions
99     validMask(validMask == -1) = 0;
100
101     % Get a patch of windowSize x windowSize
        around the pixel whose value we need to
        find
102     patch1 = im_synth(max(1, pixelList(p, 1) -
        halfWinSize): min(outSize, pixelList(p,

```

```

, 1) + halfWinSize), max(1, pixelList(p
, 2) - halfWinSize): min(outSize,
pixelList(p, 2) + halfWinSize)).*
validMask;

103
104 %% For every patch in the input image,
    compare it with the patch we got for
    the pixel we want to compute
105 for x = 1:w-winsize
106     for y = 1:h-winsize
107         if size(im(x:x+winsize-1,y:y+
            winsize-1)) == size(validMask)
108             % Apply Valid Mask
109             patch2 = im(x:x+winsize-1,y:y+
                winsize-1).*validMask;
110             if size(patch1) == size(patch2
                )
111                 % Find the minimum SSD
112                 dist = (patch1 - patch2);
113                 ssd = sum(dist(:).^2);
114             end
115             % Find the pixels in the input
                image have a similar
                neighbourhood w.r.t. the
                filled neighbours of the
                currently unfilled pixel in
                the output image
116             if ssd < ssd_min
117                 % Get minimum SSD
118                 ssd_min = ssd;
119                 patch3 = im(x:x+winsize-1,
                    y:y+winsize-1);
120             end
121         end
122     end
123 end

124
125 %% Get the best matches
126 bestMatches = zeros(w*h, 1);
127 count1 = 1;
128 for x = 1:w

```

```

129         for y = 1:h
130             pix = im(x, y);
131             if ssd_min*(1+error_threshold) >=
                pix
132                 bestMatches(count1, 1) = pix;
133                 count1 = count1 + 1;
134             end
135         end
136     end
137
138     %% Randomly pick an input image pixel from
        BestMatches and paste its pixel value
        into the current unfilled output pixel
        location
139     im_synth(pixelList(p, 1), pixelList(p, 2))
        = bestMatches(randperm(size(
            bestMatches, 1), 1), 1);
140
141     end
142 end
143 end

```

C

Matlab code for Extension for Extra Credit

C.1 Implementation of synthImageQuilting.m

```
1 function im_patch = synthImageQuilting(im, tileSize ,
    numTiles, outSize)
2     %% Get the dimensions of the input image
3     [w, h, c] = size(im);
4
5     %% Initialize the output image with all 0s
6     im_patch = zeros(outSize);
7
8     %% Get random tile co-ordinate pairs for the seed
    patch
9     x = floor(1 + rand(1,1) * (w - tileSize - 1));
10    y = floor(1 + rand(1,1) * (h - tileSize - 1));
11
12    % Get the seed patch
13    patch = im(x:x+tileSize-1, y:y+tileSize-1);
14    % Copy the random tile seed patch from source
    image to the top-left corner
15    im_patch(1:tileSize, 1:tileSize) = patch;
16
17    %% size of the region of overlap between tiles –
    typically 1/6
18    overlap = ceil(1/6 * tileSize);
19
20    %% For each tile of the output/synthesized image
21    for i = 1:numTiles
22        for j = 1:numTiles
23
24            % Do nothing for the first tile as it is
                already set by the random patch
25            if (i == 1 && j == 1)
26                continue;
27            end
28
29            ssd_C_min = inf;
30            ssd_D_min = inf;
```

```

31
32      %% SSD (sum of squared difference) between
           the neighbors of the output image tile
           position and each tile position in the
           source image, considering the region
           of overlap
33      for m = 1 : w - tileSize
34          for n = 1 : h - tileSize
35              temp_patch = im(m:m+tileSize-1, n:
                           n+tileSize-1);
36
37              % Upper Horizontal overlapping
                 section
38              A = temp_patch(1:tileSize, 1:
                           overlap);
39              % Left Vertical overlapping
                 section
40              B = temp_patch(1:overlap, 1:
                           tileSize);
41
42              if (i ~= 1 && j ~= 1)
43                  C = im_patch((i-1)*tileSize
                               +1:(i-1)*tileSize+tileSize,
                               (j-1)*tileSize-overlap:(j
                               -1)*tileSize-1);
44                  diff = A - C;
45                  ssd_C = sum(sum(diff(:).^2));
46                  D = im_patch((i-1)*tileSize-
                               overlap:(i-1)*tileSize-1, (
                               j-1)*tileSize+1:(j-1)*
                               tileSize+tileSize);
47                  diff = B - D;
48                  ssd_D = sum(sum(diff(:).^2));
49                  if ssd_C + ssd_D < ssd_C_min +
                     ssd_D_min
50                      patch = temp_patch;
51                      ssd_C_min = ssd_C;
52                      ssd_D_min = ssd_D;
53                  end
54              elseif (i ~= 1)
55                  D = im_patch((i-1)*tileSize-

```



```

        overlap:(i-1)*tileSize-1, (
        j-1)*tileSize+1:(j-1)*
        tileSize+tileSize);
56     diff = B - D;
57     ssd_D = sum(sum(diff(:).^2));
58
59     if ssd_D < ssd_D_min
60         patch = temp_patch;
61         ssd_D_min = ssd_D;
62     end
63 elseif(j ~= 1)
64     C = im_patch((i-1)*tileSize
        +1:(i-1)*tileSize+tileSize,
        (j-1)*tileSize-overlap:(j
        -1)*tileSize-1);
65     diff = A - C;
66     ssd_C = sum(sum(diff(:).^2));
67
68     if ssd_C < ssd_C_min
69         patch = temp_patch;
70         ssd_C_min = ssd_C;
71     end
72 end
73
74     end
75 end
76
77 %% Copy the tile in source image with
    lowest SSD (minSSD) with respect to the
    current tile in output image, to the
    output image
78 im_patch((i-1)*tileSize+1:(i-1)*tileSize+
    tileSize, (j-1)*tileSize+1:(j-1)*
    tileSize+tileSize) = patch;
79
80 %% To the next un-filled position in the
    output image (left-to-right, top-to-
    bottom)
81     end
82 end
83

```

84 end